

Database Fundamentals

NashTech Viet Nam

09th Apr 2020




Warm up - Introductions

- Your role
- Your background and experience in the subject

Course Audience and Prerequisite

- Audience:
 - Fresh Graduate, Junior Engineer
- Prerequisite
 - None

Agenda

- 
- Module 1: Introducing Core Database Concepts
 - Module 2: Understanding Data Modeling
 - Module 3: Creating Database and Objects
 - Module 4: Manipulating and Querying Data
 - Module 5: Administering Database

Introducing Core Database Concepts



Module Overview

- Databases Concepts
- Normalization
- Referential integrity
- Constraints

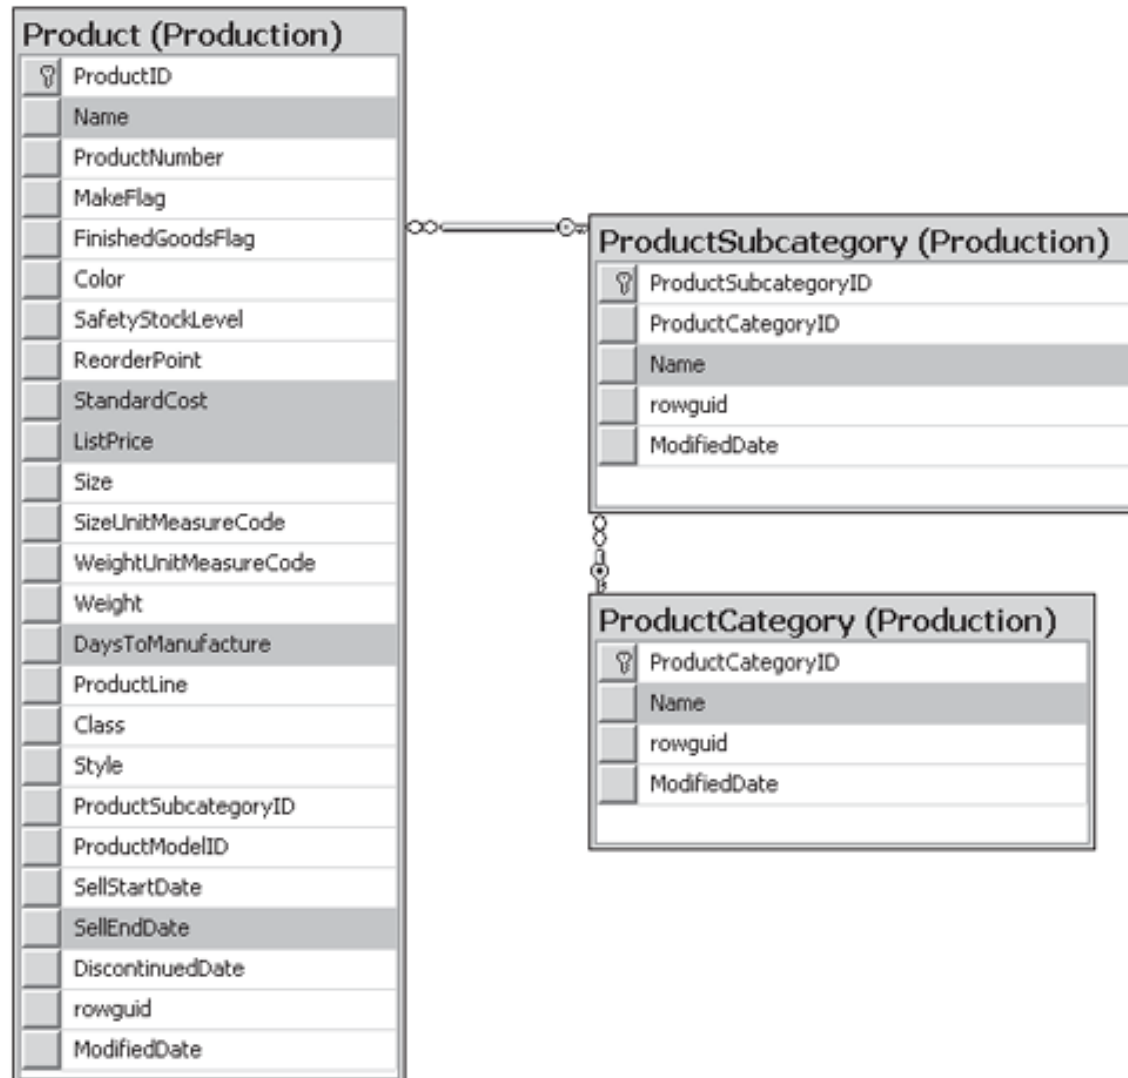
Database

- A ***database (db)*** is an organized collection of data, typically stored in electronic format
 - It allows you to input, manage, organize, and retrieve data quickly
 - Traditional databases are organized by records (rows), fields (columns) stored in tables which are stored in the database files

Relational databases

- A ***relational database*** a collection of tables of data all of which are formally described and organized according to the relational model. Each table must identify a column or group of columns, called the ***PRIMARY KEY***, to uniquely identify each row

Sample relational structure



Referential integrity

- ***Referential Integrity*** (RI) is a database concept used to ensure that the relationships between your database tables remains synchronized during data modifications.
- RI can be used to ensure the data is clean, may be helpful in optimizing your database environment and can assist in early detection of errors.
- A combination of **PRIMARY KEY** and **FOREIGN KEY** constraints can be used to help enforce referential integrity of your database.

Database Management System (DBMS)

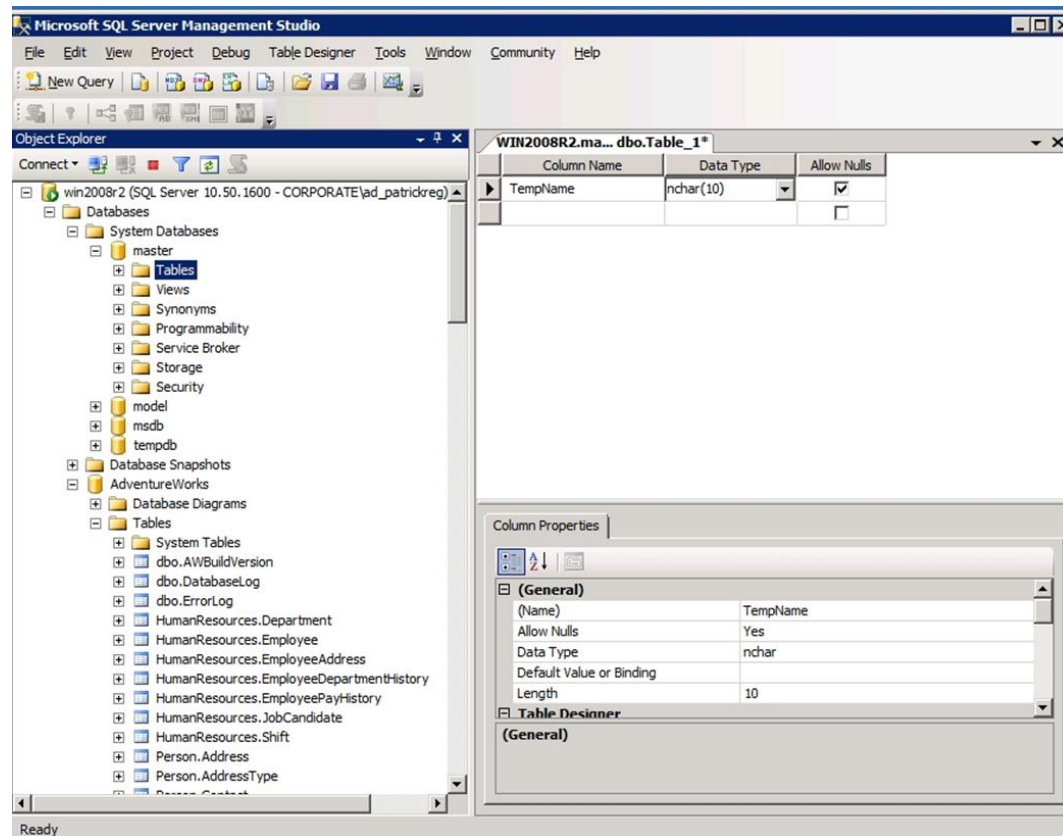
- **Database Management System (*DBMS*)** is a software system designed to allow the definition, creation, querying, and updating of data stored in databases.
- ***DBMS*** is used by the users to access the data stored in database files. A DBMS is also used to perform administrative tasks on the databases and objects contained within the database.
- A **RDBMS** also provides what DBMS provides but above that it provides relationship integrity.
- In short, we can say **RDBMS = DBMS + Referential Integrity**

Database servers

- Databases are stored on **database servers** which are dedicated physical or virtual servers that host the database files and provide high-level performance for users who are accessing the data.
- Database servers contain the DBMS used to manage the data and administer the SQL Server environment.
- A database server can have one default instance and several named instances of SQL Server. A SQL Server instance is a copy of the sqlservr.exe program that runs as a Windows operating system service.
- Often multiple database servers are deployed to provide high availability and improve performance

SQL Server Management Studio (SSMS)

- A graphical user interface (GUI) used to browse, select, and manage the SQL Server instance and any of the objects within that SQL Server instance.



Normalizing a database

- Normalization the process of organizing data in a database that includes creating tables and establishing relationships between the tables
- Process is used to help eliminate redundant data
- Three typical ***normalization forms*** (NFs)
 - 1NF: Eliminate Repeating Groups
 - 2NF: Eliminate Redundant Data
 - 3NF: Eliminate Columns Not Dependent on Key

First normal form (1NF)

- A database is in first normal form if it satisfies the following conditions:
 - Contains only atomic values
 - There are no repeating groups
- Do not use multiple fields in a single table to store similar data

Example of 1ST Normal Form(1NF)

- Un-normalized table:

TABLE_PRODUCT

Product ID	Color	Price
1	red, green	15.99
2	yellow	23.99
3	green	17.50
4	yellow, blue	9.99
5	red	29.99

- 1ST Normal Form: No repeating group



TABLE_PRODUCT_PRICE

Product ID	Price
1	15.99
2	23.99
3	17.50
4	9.99
5	29.99

TABLE_PRODUCT_COLOR

Product ID	Color
1	red
1	green
2	yellow
3	green
4	yellow
4	blue
5	red

Example of normalization

■ Un-normalized table

Student#	Advisor	Adv-Room	Class1	Class2	Class3
1022	Jones	412	101-07	143-01	159-02
4123	Smith	216	201-01	211-02	214-01

■ First Normal Form: No Repeating Groups

Student#	Advisor	Adv-Room	Class#
1022	Jones	412	101-07
1022	Jones	412	143-01
1022	Jones	412	159-02
4123	Smith	216	201-01
4123	Smith	216	211-02
4123	Smith	216	214-01

Second normal form (2NF)

- A database is in second normal form if it satisfies the following conditions:
 - It is in first normal form
 - All non-key attributes are fully functional dependent on the primary key

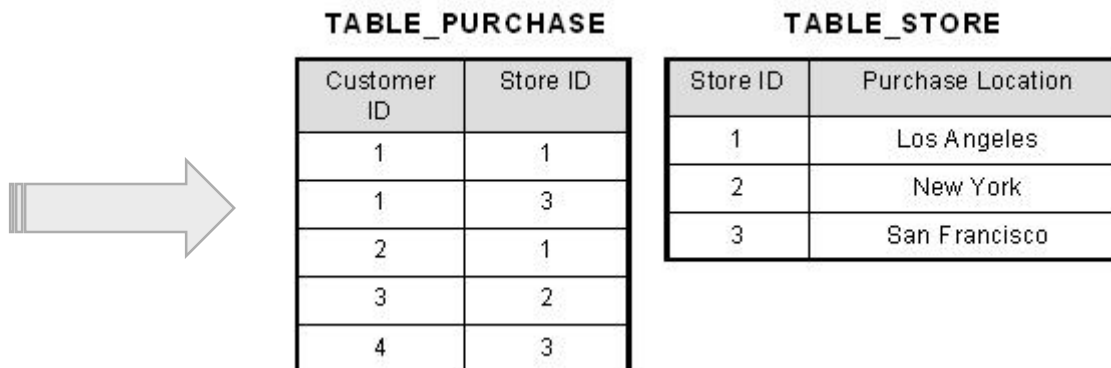
Example of 2nd Normal Form(2NF)

- Un-normalized table:

TABLE_PURCHASE_DETAIL

Customer ID	Store ID	Purchase Location
1	1	Los Angeles
1	3	San Francisco
2	1	Los Angeles
3	2	New York
4	3	San Francisco

- 2nd Normal Form: eliminate redundant data



Third normal form (3NF)

- A database is in third normal form if it satisfies the following conditions:
 - It is in second normal form
 - There is no transitive functional dependency

Example of 3rd Normal Form(3NF)

- Un-normalized table:

TABLE_BOOK_DETAIL

Book ID	Genre ID	Genre Type	Price
1	1	Gardening	25.99
2	2	Sports	14.99
3	1	Gardening	10.00
4	3	Travel	12.99
5	2	Sports	17.99

- 3rd Normal Form: Third Normal Form: eliminate data not dependent on the key



TABLE_BOOK

Book ID	Genre ID	Price
1	1	25.99
2	2	14.99
3	1	10.00
4	3	12.99
5	2	17.99

TABLE_GENRE

Genre ID	Genre Type
1	Gardening
2	Sports
3	Travel

Methods for enforcing referential integrity

- There are several methods available in SQL Server to help maintain database integrity:
 - Primary key constraint
 - Foreign key constraint
 - Unique constraint
 - Indexes
 - Triggers
- Any of these methods can be created as a ***composite key*** which is an index or constraint created using more than one column.

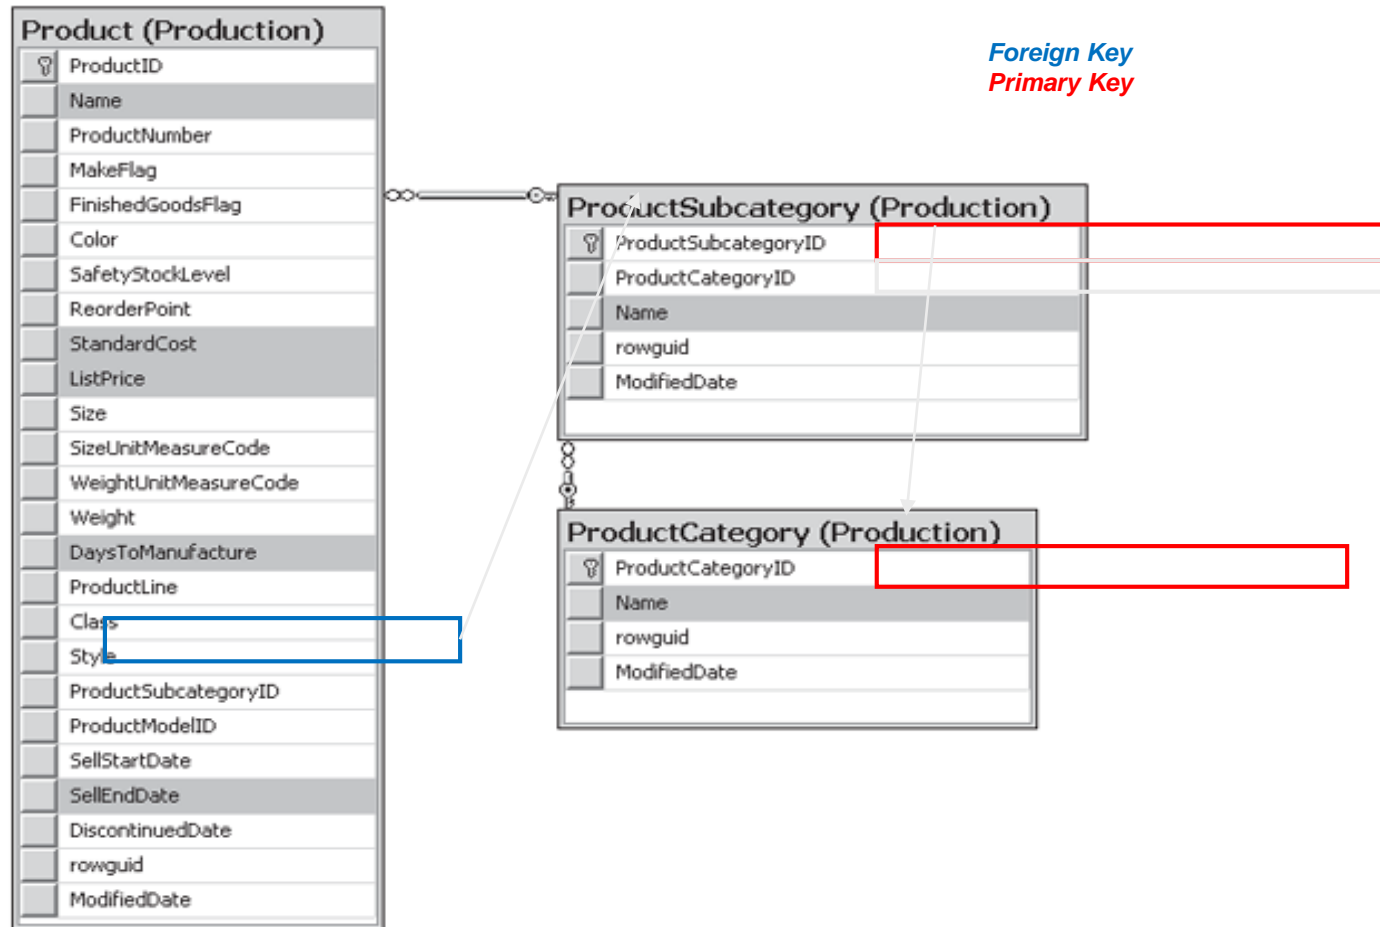
PRIMARY KEY constraint

- An important concept of designing a database table is the use of a ***PRIMARY KEY*** — an attribute or set of attributes used to uniquely identify each row
- A table can only have one primary key which is created using a primary key constraint and enforced by creating a unique index on the primary key columns
- A column that participates in the primary key constraint cannot accept null values

FOREIGN KEY constraint

- A ***FOREIGN KEY*** is a column or combination of columns that are used to establish a link between data in two tables.
- A foreign key does not have to reference a primary key, it can be defined to reference a unique constraint in either the same table or in another table

Relational structure with keys



Agenda

- Module 1: Introducing Core Database Concepts
- ➡ ■ Module 2: Understanding Data Modeling
- Module 3: Creating Database and Objects
- Module 4: Manipulating and Querying Data
- Module 5: Administering Database

Module Overview

- Data Modeling
- Data Model
- Different Data Models
 - Conceptual Data Model
 - Logical Data Model
 - Physical Data Model

Data Modeling

- Process of creating data model for an Information System by applying formal data modeling techniques
- Process used to define and analyze data requirement needed to support business process
- Who involves
 - Data modelers
 - Business Stakeholders
 - Potential users of Information System

What is Data Model?

- Data Model is a collection of conceptual tool for describing data, data relationship, data semantics and consistency constraint

Different Data Model


- Conceptual: defines WHAT the system contains
- Logical: describe HOW the system will be implemented, regardless of DBMS
- Physical: describe HOW the system will be implemented using a specific DBMS
- Data Model elements
 - Entity: a real world thing or an interaction between 2 or more real world things
 - Attribute: an atomic piece of information that we need to know about Entity
 - Relationship: How entity depend on each other
 - ONE-TO-ONE
 - ONE-TO-MANY
 - MANY-TO-MANY

Different features in Data Model

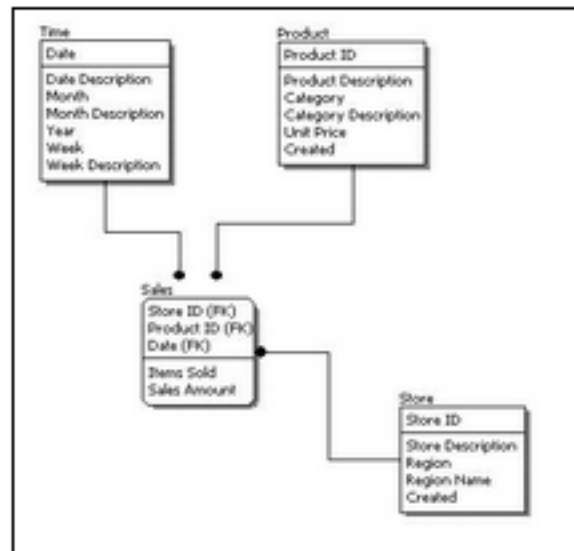
Feature	Conceptual	Logical	Physical
Entity Names	✓	✓	
Entity Relationships	✓	✓	
Attributes		✓	
Primary Keys		✓	✓
Foreign Keys		✓	✓
Table Names			✓
Column Names			✓
Column Data Types			✓

Conceptual, logical, and physical versions of a single data model

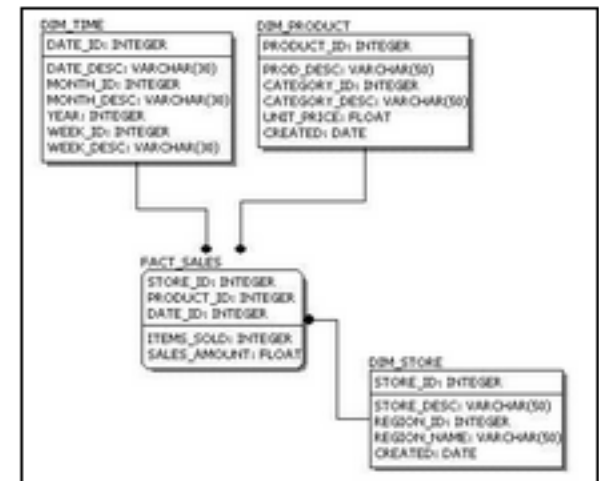
Conceptual Model Design



Logical Model Design



Physical Model Design



Agenda

- Module 1: Introducing Core Database Concepts
- Module 2: Understanding Data Modeling
- Module 3: Creating Database and Objects
- Module 4: Manipulating and Querying Data
- Module 5: Administering Database



Module Overview

- Data types
- Database objects
- DDL statements

Data types

- A ***data type*** is an attribute that specifies the type of data that an object can hold as well as the number of bytes of information that can be stored in the object
- If you have similar data types to choose from but they only differ in byte size, use the data type that has a larger range of values and/or has increased precision
- Exact numeric data types (int, tinyint) are the most common SQL Server data types used to store numeric information.
- Approximate Numerics include precision (p) which is the total number of decimal digits that could be stored, both to the left and right of the decimal point.

Built-in data type categories

- SQL Server built-in data types are organized into the following categories:
 - Exact numerics – (bigint, bit, decimal, int, money, numeric, smallint)
 - Approximate numerics (float, real)
 - Date and time (date, datetime2, datetime, datetimeoffset, time)
 - Character strings (char, varchar, text)
 - Unicode character strings (nchar, ntext, nvarchar)
 - Binary strings (binary, varbinary, image)
 - Other data types (cursor, timestamp, uniqueidentifier, table)
 - Large valued data types (varchar(max), nvarchar(max))
 - Large object data types (text, ntext, image, xml)

Data types

- **Money** - used where you'll store money or currency values
- **Int** - used to store whole numbers and when performing mathematical computations
- **Float** - commonly used in the scientific community and is considered an approximate-number data type
- **Datetime** - used to store date and time values in one of many different formats

Data types

- **Char** – fixed length non-unicode string data type where n defines the string length
- **Varchar** – variable length non-unicode string data type that indicates the actual storage size of the data
- **Bit** (Boolean) – integer that can have a null, 0 (False), or 1 (True) value

Data types storage size

Data Type	Use/Description	Storage Size
Int	Integer data from -2^{31} (-2,147,483,648) to $2^{31}-1$ (2,147,483,647)	4 bytes
Float	Approximate number - $1.79E+308$ to $-2.23E-308$, 0 and $2.23E-308$ to $1.79E+308$	Depends on the value of n
Datetime	Date Range January 1, 1753, through December 31, 9999 Time Range 00:00:00 through 23:59:59.997	8 bytes
Char	Fixed-length, non-Unicode string data. Can be a value from 1 through 8,000	Defined width
Varchar	Variable-length non-Unicode string. Can be a value from 1 through 8,000	Actual length + 2 bytes
Bit	Integer with a value of 0 or 1.	1 byte for every 8 bit columns

Implicit and explicit conversions

- Implicit data type conversions occurs when the SQL Server expression evaluator automatically converts data from one data type to another to complete an operation like a comparison of two values
- Explicit data type conversions require the use of the CONVERT or CAST function to convert data from one data type to another before an operation like a comparison can be completed.
- To convert a numeric value into a character string
 - ✓ **CAST** (\$157.27 **AS VARCHAR**(10))
- Not all data types conversions are supported
 - ✓ *nchar* cannot be converted to *image*
- Use CAST instead of CONVERT to adhere to ISO
- Use CONVERT instead of CAST to take advantage of the style functionality

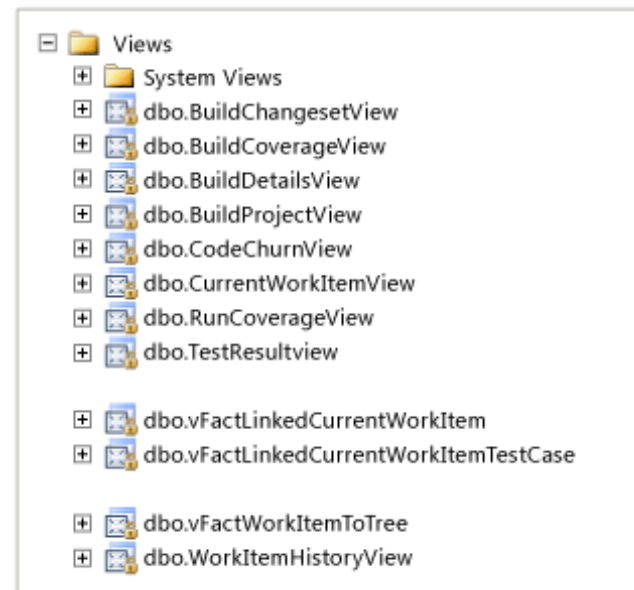
Tables

- A **table** is a collection of rows and columns that is used to organize information about a single topic. Each row within a table corresponds to a single record and contains several attributes that describe the row.

EmployeeID	LastName	FirstName	Department
100	Smith	Bob	IT
101	Jones	Susan	Marketing
102	Adams	John	Finance

Views

- A **view** is simply a virtual table consisting of different columns from one or more tables.
- Unlike a table, a view is stored in the database as a query object; therefore, a view is an object that obtains its data from one or more underlying tables.



Stored procedures

- A ***stored procedure*** is a group of Transact-SQL statements that have been compiled and saved so it can be run several times.
- Parameters can be passed to and returned from a stored procedure so they can be reused with different values.

```
IF (@QuantityOrdered < (SELECT QuantityOnHand
                        FROM Inventory
                        WHERE PartID = @PartOrdered) )
    BEGIN
        -- SQL statements to update tables and process order.
    END
ELSE
    BEGIN
        -- SELECT statement to retrieve the IDs of alternate items
        -- to suggest as replacements to the customer.
    END
```

User-Defined functions

- ***User-defined functions (udf)*** are routines that takes zero or more parameters, completes an operation, and return the result of the operation as a value.
- There are three types of functions
 - ✓ Scalar – returns a single data value
 - ✓ Table-valued – returns a table data type
 - ✓ System – Provided by SQL Server, cannot be modified

Primary differences between SP and UDF

- Stored Procedures(SP)
 - ✓ Called independently using EXEC statement
 - ✓ Cannot JOIN stored procedures
 - ✓ Can be used to modify SQL Server configuration
 - ✓ Can use nondeterministic functions such as GETDATE()

- User-defined Functions(UDF)
 - ✓ Called from within another SQL statement
 - ✓ Can JOIN UDF's
 - ✓ Cannot be used to modify SQL Server configuration
 - ✓ Always stops execution of T-SQL code if error occurs

Common DDL statements

- CREATE – define new entities
- ALTER – modify existing entities
- DROP – remove existing entities

CREATE statement

- Used to create new entities in SQL Server including some of the most common entities

✓ Database	Procedure
✓ Table	Trigger
✓ Default	View
✓ Index	User
✓ Login	Role

```
CREATE DATABASE Sales ON ( NAME = Sales_dat, FILENAME =  
'C:\Program Files\Microsoft SQL  
Server\MSSQL11.MSSQLSERVER\MSSQL\DATA\sales.mdf', SIZE =  
10, MAXSIZE = 50, FILEGROWTH = 5 )  
LOG ON ( NAME = Sales_log, FILENAME = 'C:\Program Files\Microsoft  
SQL Server\MSSQL11.MSSQLSERVER\MSSQL\DATA\salelog.ldf', SIZE  
= 5MB, MAXSIZE = 25MB, FILEGROWTH = 5MB ) ;
```

Create new table

- USE SALES
- GO
- --Create new table called Products
CREATE TABLE dbo.Products1
(
 ProductID int NULL,
 ProductName varchar(20) NULL,
 UnitPrice money NULL,
 ProductDescription varchar(50) NULL
);

ALTER statement

- Used to modify existing entities in SQL Server including
 - ✓ Database Trigger
 - ✓ Table View
 - ✓ Index User
 - ✓ Login Role
 - ✓ Procedure Schema

- ALTER DATABASE Sales
- Modify Name = SalesForecast ;

DROP statement

- Used to delete existing entities in SQL Server including
 - ✓ Database Trigger
 - ✓ Table View
 - ✓ Index User
 - ✓ Login Role
 - ✓ Procedure Schema
- DROP DATABASE SalesForecast

Agenda

- Module 1: Introducing Core Database Concepts
- Module 2: Understanding Data Modeling
- Module 3: Creating Database and Objects
- ➡ ■ Module 4: Manipulating and Querying Data
- Module 5: Administering Database

Module Overview

- Introducing DML statements
- Using the SELECT statement
- Modifying data using DML statements

Common DML statements

- SELECT – retrieve data
- INSERT – add data
- UPDATE – modify data
- DELETE – remove data
- BULK INSERT – Import a data file

Using the basic **SELECT** statement

- The ***SELECT*** statement is used to retrieve rows and columns from a table

SELECT * FROM tablename

- The **SELECT** statement requires the name of the table and either the * (retrieves all columns) or specific column names
- To limit the number of rows returned you can include the **WHERE** clause in the **SELECT** statement

Sample SELECT statement

```
SELECT BusinessEntityID, JobTitle, Gender  
FROM HumanResources.Employee  
WHERE BusinessEntityID <= 50
```

- Returns the following results:

BusinessEntityID	Title	Gender
-----	-----	-----
1	Chief Executive Officer	M
2	Vice President of Engineering	F
3	Engineering Manager	M
4	Senior Tool Designer	M

Multiple WHERE clauses

- You can combine several WHERE clauses in one query statement to create more specific queries.

```
SELECT BusinessEntityID, Jobtitle, VacationHours  
FROM HumanResources.Employee  
WHERE JobTitle = 'Design Engineer' AND gender = 'F' AND  
HireDate >= '2000-JAN-01'
```

```
SELECT BusinessEntityID, Jobtitle, VacationHours  
FROM HumanResources.Employee  
WHERE VacationHours > 80 OR BusinessEntityID <= 50
```


Using the BETWEEN clause

- Retrieving rows within a date range using the BETWEEN clause

```
SELECT BusinessEntityID, Jobtitle, VacationHours  
FROM HumanResources.Employee  
WHERE VacationHours BETWEEN 75 AND 100
```

Sorting the result set using ORDER By

- Sorting the result set by using the ORDER BY to specify what field to sort by.

```
SELECT BusinessEntityID, Jobtitle, VacationHours  
FROM HumanResources.Employee  
WHERE VacationHours BETWEEN 75 AND 100  
ORDER BY VacationHours
```

- You can sort in descending order by using the DESC clause.

```
SELECT BusinessEntityID, Jobtitle, VacationHours  
FROM HumanResources.Employee  
WHERE VacationHours BETWEEN 75 AND 100  
ORDER BY VacationHours DESC
```

Using the NOT clause

- Write a query to return data that specifies what you don't want returned

```
SELECT BusinessEntityID, Jobtitle, Gender  
FROM HumanResources.Employee  
WHERE NOT Gender = 'M'
```

UNION clause

- The **UNION** clause allows you to combine the rows returned from multiple SELECT statements into a single result set

```
SELECT BusinessEntityID, Jobtitle, HireDate  
FROM HumanResources.Employee  
WHERE JobTitle = 'Design Engineer'
```

UNION

```
SELECT BusinessEntityID, Jobtitle, HireDate  
FROM HumanResources.Employee  
WHERE HireDate BETWEEN '2005-01-01' AND '2005-12-31'
```

EXCEPT and INTERSECT clauses

- The **EXCEPT** clause returns distinct values from the left query that are not found on the right query

```
SELECT ProductID  
FROM Production.Product  
EXCEPT  
SELECT ProductID  
FROM Production.WorkOrder ;
```

- The **INTERSECT** clause returns any distinct values returned by both the query on the left and right sides of intersect operand

```
SELECT ProductID  
FROM Production.Product  
INTERSECT  
SELECT ProductID  
FROM Production.WorkOrder ;
```

JOIN clause

- The **JOIN** clause allows you to combine related data from multiple tables into one result set
 - **INNER JOINS** uses a comparison operator to match rows from two tables based on values in a common column that exists in both tables
 - **OUTER JOINS** (left, right, or full) includes rows from one or both tables even if they don't have matching values
 - **CROSS JOINS** return all rows from the left table with all rows from the right table. WHERE conditions should always be included.

Aggregate sample

- SQL Server provides **aggregate functions** to assist with the summarization of large volumes of data

```
SELECT COUNT (DISTINCT SalesOrderID) AS UniqueOrders,  
        AVG(UnitPrice) AS Avg_UnitPrice,  
        MIN(OrderQty)AS Min_OrderQty,  
        MAX(LineTotal) AS Max_LineTotal  
FROM Sales.SalesOrderDetail;
```

Inserting data

- You can add a new row to a table using the INSERT statement

```
INSERT INTO Production.UnitMeasure  
VALUES (N'FT', N'Feet', '20080414')
```

- You can add multiple rows to a table using the following INSERT statement

```
INSERT INTO Production.UnitMeasure  
VALUES (N'FT2', N'Square Feet ', '20080923'),  
        (N'Y', N'Yards', '20080923'),  
        (N'Y3', N'Cubic Yards', '20080923')
```

- **BULK INSERT** can be used to import a data file into a table with a user-specified format.

Update statement

- The ***UPDATE*** statement is used to modify the data that is already stored in a table

```
UPDATE Sales.SalesPerson  
SET Bonus = 6000, CommissionPct = .10, SalesQuota = NULL  
WHERE sales.SalesPerson.BusinessEntityID = 289
```

DELETE statement

- The ***DELETE*** statement is used to delete rows from a table

```
DELETE FROM Production.UnitMeasure  
WHERE Production.UnitMeasure.Name = 'Feet'
```

- A **DELETE** statement without a WHERE clause will cause all rows to be deleted

```
DELETE FROM Sales.SalesPersonQuotaHistory;
```

Agenda

- Module 1: Introducing Core Database Concepts
- Module 2: Understanding Data Modeling
- Module 3: Creating Database and Objects
- Module 4: Manipulating and Querying Data
- Module 5: Administering Database



Module Overview

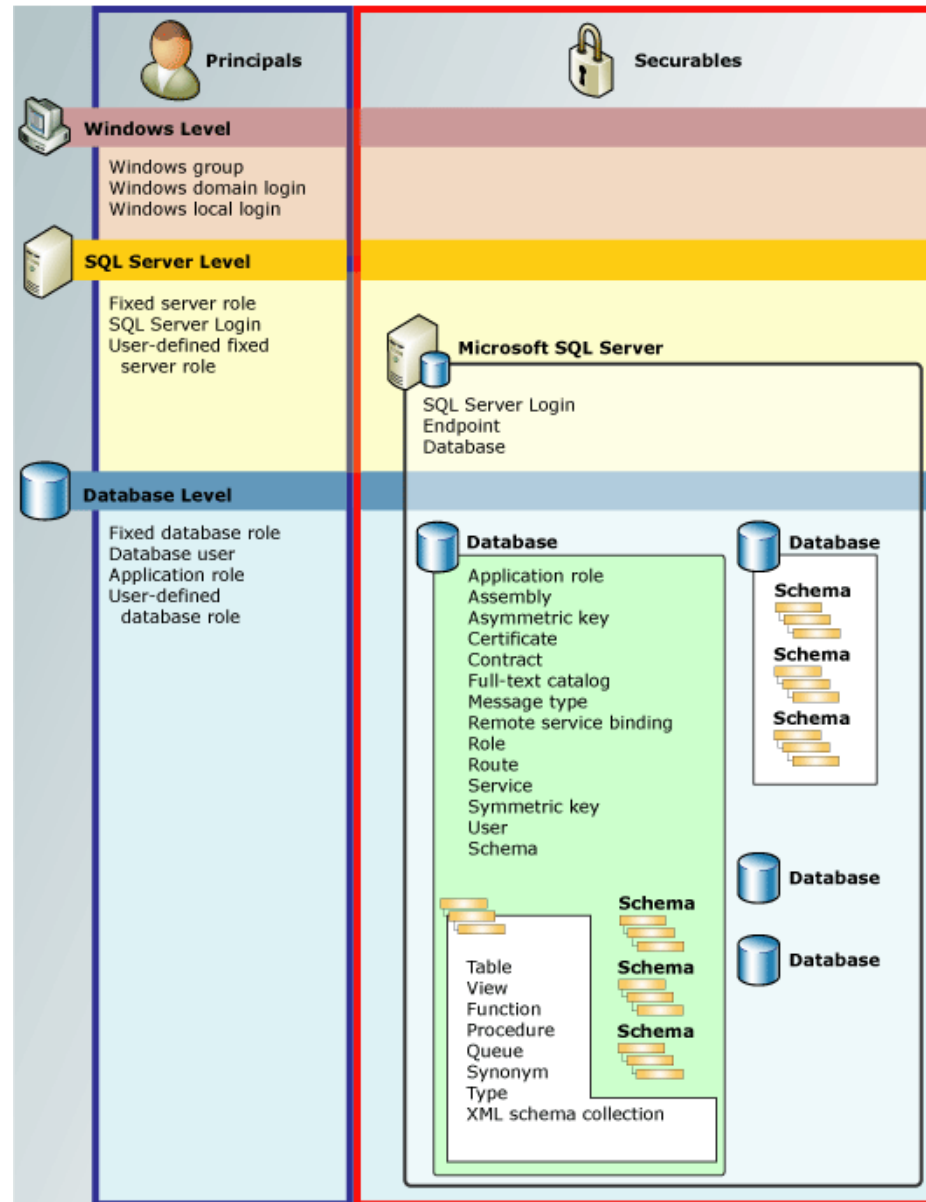
- Understanding SQL Server Security
- Securing SQL Server databases and objects

Database security

Securing your database content is a critical part of a DBA's job. The design, testing, and implementation of security is necessary to ensure that confidentiality is not compromised

- **Securables** are the server, database, and objects a database contains
- **Principals** are the individuals, groups, and processes granted access to SQL Server
- **Permissions** are granted to a principal for every SQL Server securable

Permissions Hierarchy



Logins and accounts

- Three tiered approach to accessing content
 1. SQL Server access - a **login** is a security principal that can be authenticated by a secure system to provide a user access to SQL Server
 2. Database access - a **database user** is mapped to a SQL login and provides a user or group access to a database
 3. Object access – **permissions** are applied at the object level to provide the appropriate access to the objects within the database

Server-level security

- ***Authentication*** is the act of verifying a user or system identity and allowing them to login using:
 - Windows Authentication
 - Mixed-Mode (Windows and SQL logins)
- Logins can be populated into the fixed server roles or in user-defined server roles

Fixed server roles

- SQL Server includes several fixed server roles:
 - **Sysadmin** – perform any activity on the server
 - **Dbcreator** – create, alter, drop, restore databases
 - **Securityadmin** –manage logins and their properties
- You can also create user-defined server roles that have specific permissions applied to the roles

Database-level security

- A database user is a database level security principal that must be mapped to a login at the server level in order for the user to connect to the database
- A login can be mapped to different databases as different users but can only be mapped as one user in each database
- Database users can be populated into the fixed database roles or in a user-defined database role
- All users are automatically members of the *public* database role and cannot be removed

Fixed database roles

- SQL Server includes several fixed database roles
 - db_owner: perform all configuration activities
 - db_datareader: read all data from all user tables
 - db_datawriter: add, delete, or change data
- You can also create user-defined database roles that have specific permissions applied to the roles

Guest logon accounts

- The ***guest*** user account is included in every database and is used by any user who accesses the database but does not have a user account within the database
- The guest user account cannot be dropped but it can be disabled by revoking its connect permission
REVOKE connect FROM guest

Managing object permissions

- Permissions to an object can be managed by using the following commands
 - ***Grant*** - provides a level of access to the object
 - ***Deny*** - overrides any grant permission
 - ***Revoke*** - removes the previously assigned permission, regardless of whether it was a deny or grant permission

Object permissions

- ***Object permissions*** are the permissions that allow a user to perform actions on database objects (such as tables, stored procedures, and views):
 - SELECT
 - INSERT
 - UPDATE
 - DELETE
 - EXECUTE (stored procedures)

A vibrant blue liquid splash is captured in mid-air, creating a complex, organic shape with multiple peaks and valleys. The splash is set against a background that transitions from a deep blue on the left to a bright red on the right. A white rectangular frame is superimposed over the center of the splash, containing the text "Thank you".

Thank you