

PostgreSQL Fundamental

- Long Hoang
- 2023

Nash
Tech.

Agenda

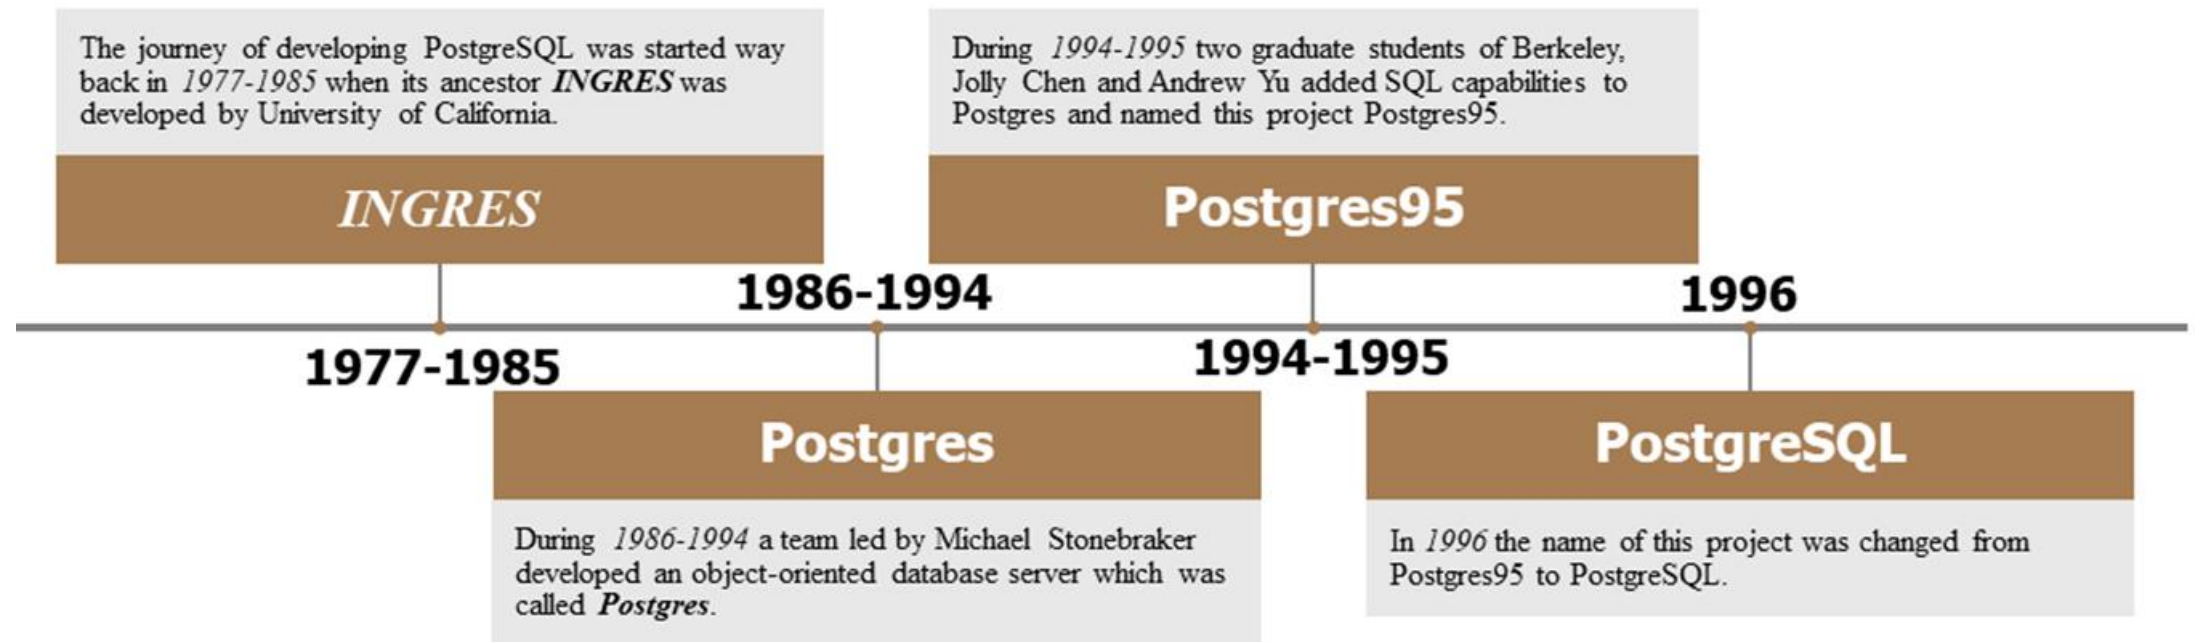
Session 1

- Introduction to PostgreSQL
- Data modeling

Session 2

- PostgreSQL SQL language

PostgreSQL Timeline

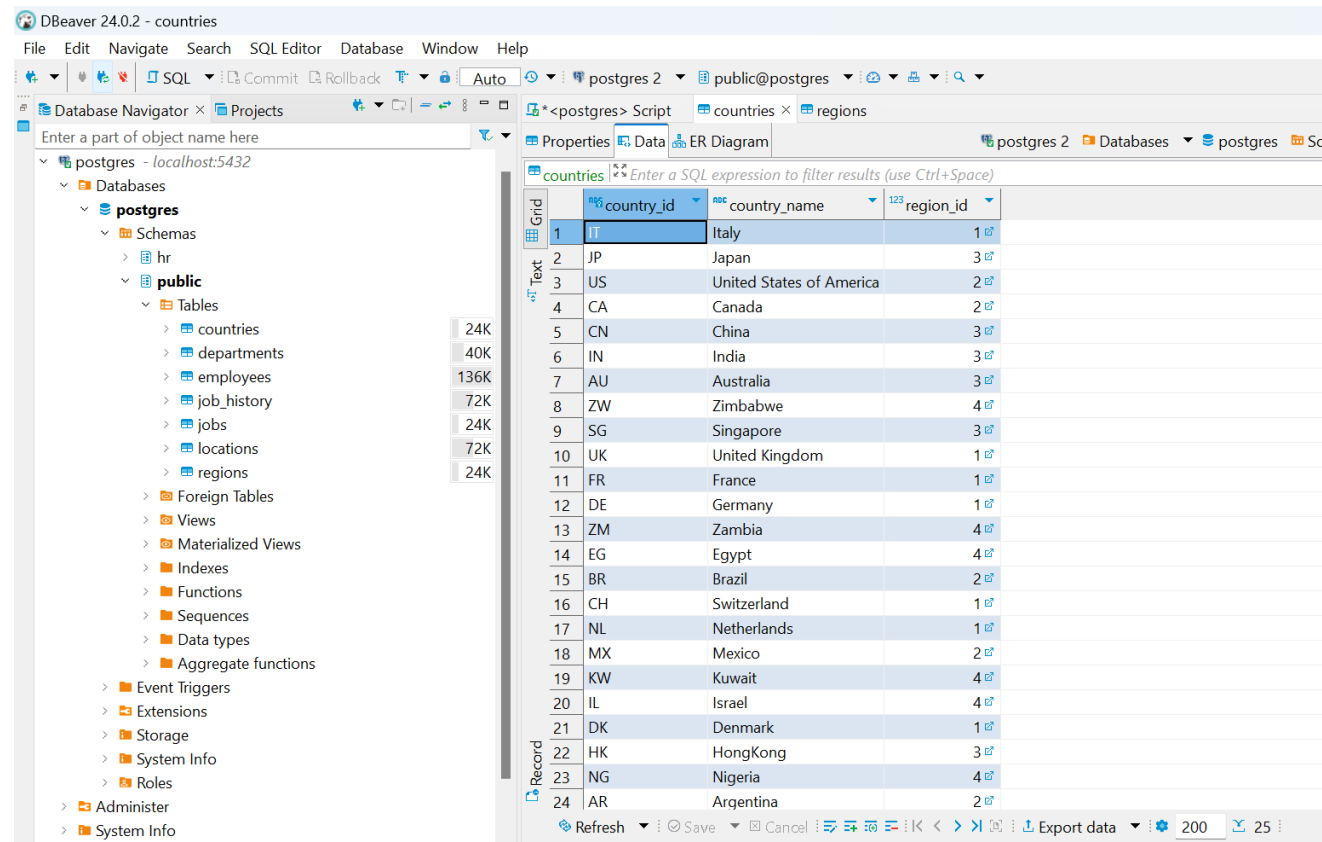


What is PostgreSQL?

- It is a relational database management system.
- It is a free and open source
- It supports almost all available operating systems.


Tools to connect database

- Command line (CMD)
- IDE like PGAdmin or DBeaver community edition
- Dbeaver is recommended for its rich features (schema compare, ERD generation, DDL/DML)



Client tool

- DBeaver community



Data Modeling

- Process of creating data model for an Information System by applying formal data modeling techniques
- Process used to define and analyze data requirement needed to support business process
- Who involves
 - Data modelers
 - Business Stakeholders
 - Potential users of Information System

What is Data Model?

- Data Model is a collection of conceptual tool for describing data, data relationship, data semantics and consistency constraint

Type of Data Model


- Conceptual: defines WHAT the system contains
- Logical: describe HOW the system will be implemented, regardless of DBMS
- Physical: describe HOW the system will be implemented using a specific DBMS
- Data Model elements
 - Entity: a real-world thing or an interaction between 2 or more real world things
 - Attribute: an atomic piece of information that we need to know about Entity
 - Relationship: How entity depend on each other
 - ONE-TO-ONE
 - ONE-TO-MANY
 - MANY-TO-MANY

Features in Data Modeling process


Feature	Conceptual	Logical	Physical
Entity Names	✓	✓	
Entity Relationships	✓	✓	
Attributes		✓	
Primary Keys		✓	✓
Foreign Keys		✓	✓
Table Names			✓
Column Names			✓
Column Data Types			✓

An example of model design

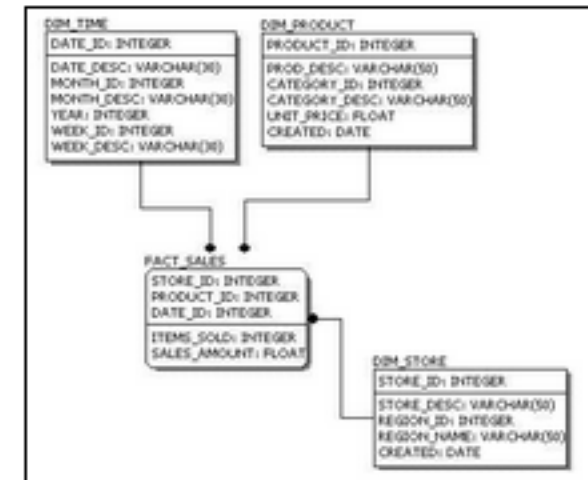
Conceptual Model Design



Logical Model Design




Physical Model Design



Session 2

Structured query language (SQL)



Module 1: Introduction to SQL

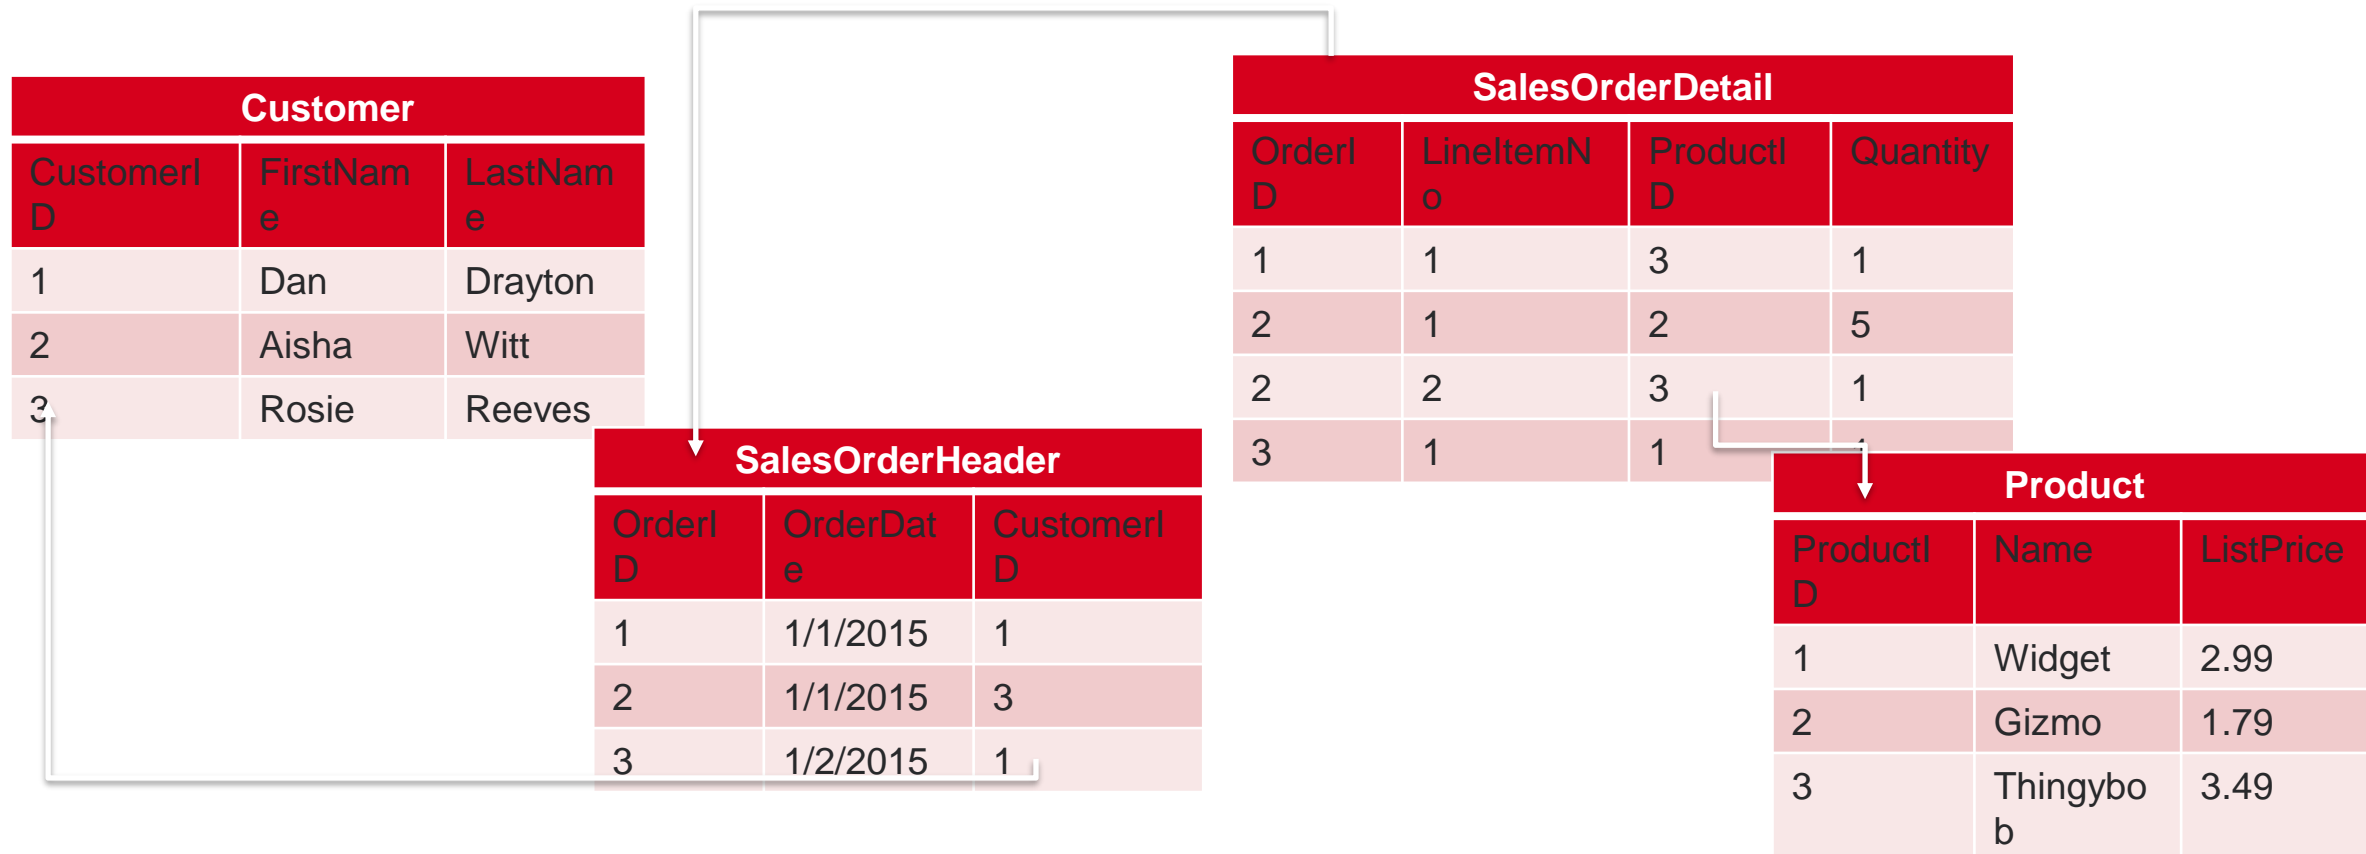
- What is SQL?
- Relational Databases
- Schemas and Object Names
- SQL Statement Types
- The SELECT Statement
- Working with Data Types
- Working with NULLs

What is SQL?

- Structured Query Language (SQL)
 - Developed by IBM in 1970s
 - Adopted as a standard by ANSI and ISO standards bodies
 - Widely used in industry
- SQL is declarative, not procedural
 - Describe what you want, don't specify steps

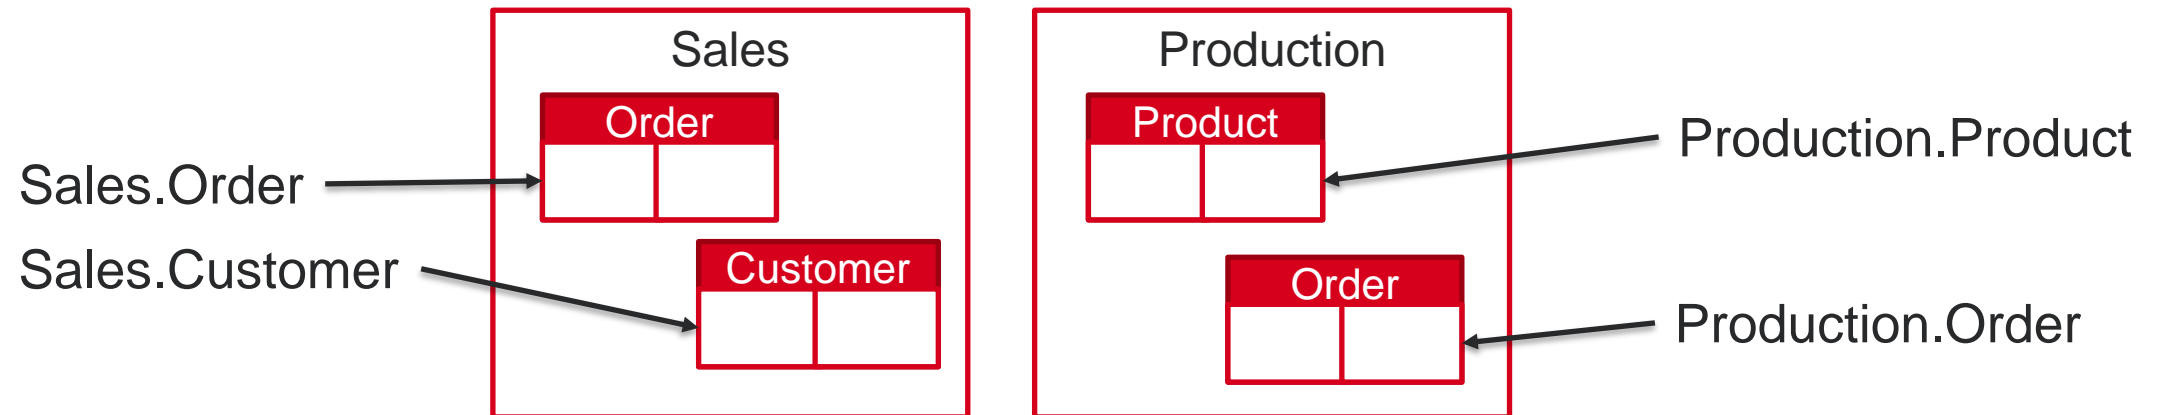
Relational Databases

- Entities are represented as *relations* (tables), in which their attributes are represented as *domains* (columns)
- Most relational databases are *normalized*, with relationships defined between tables through *primary* and *foreign* keys



Schemas and Object Names

- Schemas are namespaces for database objects
- Fully-qualified names:
`[server_name.][database_name.][schema_name.]object_name`
- Within database context, best practice is to include schema name:
`schema_name.object_name`



SQL Statement Types

Data Manipulation Language (DML)	Data Definition Language (DDL)	Data Control Language (DCL)
Statements for querying and modifying data: <ul style="list-style-type: none">• SELECT• INSERT• UPDATE• DELETE	Statements for defining database objects: <ul style="list-style-type: none">• CREATE• ALTER• DROP	Statements for assigning security permissions: <ul style="list-style-type: none">• GRANT• REVOKE• DENY



Focus of this course

The SELECT Statement

	Element	Expression	Role
5	SELECT	<select list>	Defines which columns to return
1	FROM	<table source>	Defines table(s) to query
2	WHERE	<search condition>	Filters rows using a predicate
3	GROUP BY	<group by list>	Arranges rows by groups
4	HAVING	<search condition>	Filters groups using a predicate
6	ORDER BY	<order by list>	Sorts the output

```
SELECT OrderDate, COUNT(OrderID)
FROM Sales.SalesOrder
WHERE Status = 'Shipped'
GROUP BY OrderDate
HAVING COUNT(OrderID) > 1
ORDER BY OrderDate DESC;
```

Basic SELECT Query Examples

- All columns

```
SELECT * FROM Production.Product;
```

- Specific columns

```
SELECT Name, ListPrice  
FROM Production.Product;
```

- Expressions and Aliases

```
SELECT Name AS Product, ListPrice * 0.9 AS SalePrice  
FROM Production.Product;
```

Working with Data Types

Numeric	Character	Date/Time
smallint	char(n)	date
int	varchar(n)	time
bigint	text	timestamp
decimal		
numeric		
double		
real		

Data Type Conversion

- Implicit Conversion
 - Compatible data types can be automatically converted
- Explicit Conversion
 - Requires an explicit conversion function
 - CAST

Working with NULLs – NULL values

- NULL represents a missing or unknown value
- ANSI behaviour for NULL values:
 - The result of any expression containing a NULL value is NULL
 - $2 + \text{NULL} = \text{NULL}$
 - `'MyString: ' + NULL = NULL`
 - Equality comparisons always return false for NULL values
 - $\text{NULL} = \text{NULL}$ returns *false*
 - **NULL IS NULL** returns *true*

Working with NULLs – NULL Functions

- COALESCE(*column/variable*, *value*)
 - Returns *value* if the column or variable is NULL
- WHERE column IS NULL
 - Return value is null

Module 2: Querying table with SELECT

- Removing Duplicates
- Sorting Results
- Paging Sorted Results
- Filtering and Using Predicates

Removing Duplicates

- SELECT ALL
 - Default behavior includes duplicates

```
SELECT Color
FROM Production.Product;
```

- SELECT DISTINCT
 - Removes duplicates

```
SELECT DISTINCT Color
FROM Production.Product;
```

Color
Blue
Red
Yellow
Blue
Yellow
Black

Color
Blue
Red
Yellow
Black

Sorting Results

- Use ORDER BY to sort results by one or more columns
 - Aliases created in SELECT clause are visible to ORDER BY
 - You can order by columns in the source that are not included in the SELECT clause
 - You can specify ASC or DESC (ASC is the default)

```
SELECT ProductCategory AS Category, ProductName  
FROM Production.Product  
ORDER BY Category, Price DESC;
```

Limiting Sorted Results

- LIMIT keyword allows you to limit the number of rows returned by a query
- Works with ORDER BY clause to limit rows by sort order
- Example

```
SELECT *  
FROM Employees  
ORDER BY FIRST_NAME  
LIMIT <N>
```

Eliminating Duplicates and Sorting Results

```
select first_name,last_name from public.employees
```

```
select distinct coalesce(department_id,999) as department_id from  
public.employees e order by department_id
```

```
select * from public.employees order by department_id limit 5
```

Filtering and Using Predicates

Predicates and Operators	Description
= < >	Compares values for equality / non-equality.
IN	Determines whether a specified value matches any value in a subquery or a list.
BETWEEN	Specifies an inclusive range to test.
LIKE	Determines whether a specific character string matches a specified pattern, which can include wildcards.
AND	Combines two Boolean expressions and returns TRUE only when both are TRUE.
OR	Combines two Boolean expressions and returns TRUE if either is TRUE.
NOT	Reverses the result of a search condition.

Filtering with Predicates

```
--List information about product model 6
SELECT Name, Color, Size FROM SalesLT.Product WHERE ProductModelID = 6;

--List information about products that have a product number beginning FR
SELECT productnumber, Name, ListPrice FROM SalesLT.Product WHERE ProductNumber LIKE 'FR%';

--Filter the previous query to ensure that the product number contains two sets of two digits
SELECT Name, ListPrice FROM SalesLT.Product WHERE ProductNumber LIKE 'FR-[0-9][0-9]-[0-9][0-9]';

--Find products that have no sell end date
SELECT Name FROM SalesLT.Product WHERE SellEndDate IS NOT NULL;

--Find products that have a sell end date in 2006
SELECT Name FROM SalesLT.Product WHERE SellEndDate BETWEEN '2006/1/1' AND '2006/12/31';

--Find products that have a category ID of 5, 6, or 7.
SELECT ProductCategoryID, Name, ListPrice FROM SalesLT.Product WHERE ProductCategoryID IN (5, 6, 7);

--Find products that have a category ID of 5, 6, or 7 and have a sell end date
SELECT ProductCategoryID, Name, ListPrice, SellEndDate FROM SalesLT.Product WHERE ProductCategoryID IN (5, 6, 7) AND SellEndDate IS NULL;


--Select products that have a category ID of 5, 6, or 7 and a product number that begins FR
SELECT Name, ProductCategoryID, ProductNumber FROM SalesLT.Product WHERE ProductNumber LIKE 'FR%' OR ProductCategoryID IN (5, 6, 7);
```

Module 3: Query multiple tables with JOINS

- Join Concepts
- Join Syntax
- Inner Joins
- Outer Joins
- Cross Joins
- Self Joins

Join Concepts

- Combine rows from multiple tables by specifying matching criteria
 - Usually based on primary key – foreign key relationships
 - For example, return rows that combine data from the **Employee** and **SalesOrder** tables by matching the **Employee.EmployeeID** primary key to the **SalesOrder.EmployeeID** foreign key
- It helps to think of the tables as sets in a Venn diagram



Join Syntax

- ANSI SQL-92
 - Tables joined by JOIN operator in FROM Clause
 - Preferred syntax

```
SELECT ...  
FROM   Table1 JOIN Table2  
       ON <on_predicate>;
```


- ANSI SQL-89
 - Tables joined by commas in FROM Clause
 - Not recommended: Accidental Cartesian products!

```
SELECT ...  
FROM   Table1, Table2  
WHERE  <where_predicate>;
```

Inner Joins

- Return only rows where a match is found in both input tables
- Match rows based on attributes supplied in predicate
- If join predicate operator is =, also known as equi-join

```
SELECT emp.FirstName, ord.Amount  
FROM HR.Employee AS emp  
[INNER] JOIN Sales.SalesOrder AS ord  
ON emp.EmployeeID = ord.EmployeeID
```



Using Inner Joins

-- Table aliases

```
SELECT p.Name AS ProductName, c.Name AS Category
FROM SalesLT.Product AS p
INNER JOIN SalesLT.ProductCategory AS c
ON p.ProductCategoryID = c.ProductCategoryID;
```

-- Joining more than 2 tables

```
SELECT oh.OrderDate, oh.SalesOrderNumber, p.Name AS ProductName, od.OrderQty, od.UnitPrice, od.LineTotal
FROM SalesLT.SalesOrderHeader AS oh
INNER JOIN SalesLT.SalesOrderDetail AS od
ON od.SalesOrderID = oh.SalesOrderID
INNER JOIN SalesLT.Product AS p
ON od.ProductID = p.ProductID
ORDER BY oh.OrderDate, oh.SalesOrderID, od.SalesOrderDetailID;
```


-- Multiple join predicates

```
SELECT oh.OrderDate, oh.SalesOrderNumber, p.Name AS ProductName, od.OrderQty, od.UnitPrice, od.LineTotal
FROM SalesLT.SalesOrderHeader AS oh
INNER JOIN SalesLT.SalesOrderDetail AS od
ON od.SalesOrderID = oh.SalesOrderID
INNER JOIN SalesLT.Product AS p
ON od.ProductID = p.ProductID AND od.UnitPrice = p.ListPrice --Note multiple predicates
ORDER BY oh.OrderDate, oh.SalesOrderID, od.SalesOrderDetailID;
```

Outer Joins

- Return all rows from one table and any matching rows from second table
- One table's rows are "preserved"
 - Designated with LEFT, RIGHT, FULL keyword
 - All rows from preserved table output to result set
- Matches from other table retrieved
- Additional rows added to results for non-matched rows
 - NULLs added in places where attributes do not match
- Example: Return all employees and for those who have taken orders, return the order amount. Employees without matching orders will display NULL for order amount.

```
SELECT emp.FirstName, ord.Amount  
FROM HR.Employee AS emp  
LEFT [OUTER] JOIN Sales.SalesOrder AS ord  
ON emp.EmployeeID = ord.EmployeeID;
```



Using Outer Joins

--Get all customers, with sales orders for those who've bought anything

```
SELECT c.FirstName, c.LastName, oh.SalesOrderNumber
FROM SalesLT.Customer AS c
LEFT OUTER JOIN SalesLT.SalesOrderHeader AS oh
ON c.CustomerID = oh.CustomerID
ORDER BY c.CustomerID;
```

--Return only customers who haven't purchased anything

```
SELECT c.FirstName, c.LastName, oh.SalesOrderNumber
FROM SalesLT.Customer AS c
LEFT OUTER JOIN SalesLT.SalesOrderHeader AS oh
ON c.CustomerID = oh.CustomerID
WHERE oh.SalesOrderNumber IS NULL
ORDER BY c.CustomerID;
```

--More than 2 tables

```
SELECT p.Name AS ProductName, oh.SalesOrderNumber
FROM SalesLT.Product AS p
LEFT JOIN SalesLT.SalesOrderDetail AS od
ON p.ProductID = od.ProductID
LEFT JOIN SalesLT.SalesOrderHeader AS oh --Additional tables added to the right must also use a left join
ON od.SalesOrderID = oh.SalesOrderID
ORDER BY p.ProductID;
```

Cross Joins

- Combine each row from first table with each row from second table
- All possible combinations output
- Logical foundation for inner and outer joins
 - Inner join starts with Cartesian product, adds filter
 - Outer join takes Cartesian output, filtered, adds back non-matching rows (with NULL placeholders)
- Due to Cartesian product output, not typically a desired form of join
 - Some useful exceptions:
 - Table of numbers, generating data for testing

Employee		Product	
EmployeeID	FirstName	ProductID	Name
1	Dan	1	Widget
2	Aisha	2	Gizmo

```
SELECT emp.FirstName, prd.Name
FROM HR.Employee AS emp
CROSS JOIN Production.Product AS prd;
```

Result	
FirstName	Name
Dan	Widget
Dan	Gizmo
Aisha	Widget
Aisha	Gizmo

Using Cross Joins

```
--Call each customer once per product  
SELECT p.Name, c.FirstName, c.LastName, c.Phone  
FROM SalesLT.Product as p  
CROSS JOIN SalesLT.Customer as c;
```

Self Joins

- Compare rows in same table to each other
- Create two instances of same table in FROM clause
 - At least one alias required
- Example: Return all employees and the name of the employee's manager

Employee		
EmployeeID	FirstName	ManagerID
1	Dan	NULL
2	Aisha	1
3	Rosie	1
4	Naomi	3

```
SELECT emp.FirstName AS Employee,  
       man.FirstName AS Manager  
FROM HR.Employee AS emp  
LEFT JOIN HR.Employee AS man  
ON emp.ManagerID = man.EmployeeID;
```

Result	
Employee	Manager
Dan	<i>NULL</i>
Aisha	Dan
Rosie	Dan
Naomi	Rosie

Using Self Joins

```
--note there's no employee table, so we'll create one for this example
CREATE TABLE SalesLT.Employee
(EmployeeID int IDENTITY PRIMARY KEY,
EmployeeName nvarchar(256),
ManagerID int);
GO
-- Get salesperson from Customer table and generate managers
INSERT INTO SalesLT.Employee (EmployeeName, ManagerID)
SELECT DISTINCT Salesperson, NULLIF(CAST(RIGHT(SalesPerson, 1) as INT), 0)
FROM SalesLT.Customer;
GO
UPDATE SalesLT.Employee
SET ManagerID = (SELECT MIN(EmployeeID) FROM SalesLT.Employee WHERE ManagerID IS NULL)
WHERE ManagerID IS NULL
AND EmployeeID > (SELECT MIN(EmployeeID) FROM SalesLT.Employee WHERE ManagerID IS NULL);
GO
-- Here's the actual self-join demo
SELECT e.EmployeeName, m.EmployeeName AS ManagerName
FROM SalesLT.Employee AS e
LEFT JOIN SalesLT.Employee AS m
ON e.ManagerID = m.EmployeeID
ORDER BY e.ManagerID;
```


Module 4: Using Set Operators

- What are UNION Queries?
- What are INTERSECT Queries?
- What are EXCEPT Queries?

What are UNION Queries?

- UNION returns a result set of distinct rows combined from all statements
- UNION removes duplicates during query processing (affects performance)
- UNION ALL retains duplicates during query processing

```
-- only distinct rows from both queries are returned  
SELECT countryregion, city FROM HR.Employees  
UNION  
SELECT countryregion, city FROM Sales.Customers;
```




UNION Guidelines

- Column aliases
 - Must be expressed in first query
- Number of columns
 - Must be the same
- Data types
 - Must be compatible for implicit conversion (or converted explicitly)

What are INTERSECT Queries?

- INTERSECT returns only distinct rows that appear in both result sets


```
-- only rows that exist in both queries will be returned  
SELECT countryregion, city FROM HR.Employees  
INTERSECT  
SELECT countryregion, city FROM Sales.Customers;
```



What are EXCEPT Queries?

- EXCEPT returns only distinct rows that appear in the first set but not the second
 - Order in which sets are specified matters

```
-- only rows from Employees will be returned  
SELECT countryregion, city FROM HR.Employees  
EXCEPT  
SELECT countryregion, city FROM Sales.Customers;
```



Creating UNION, INTERSECT, and EXCEPT Queries

```
-- Union example
SELECT FirstName, LastName
FROM SalesLT.Employees
UNION
SELECT FirstName, LastName
FROM SalesLT.Customers
ORDER BY LastName;
```

```
--Intersect
SELECT FirstName, LastName
FROM SalesLT.Customers
INTERSECT
SELECT FirstName, LastName
FROM SalesLT.Employees;
```

```
--Except
SELECT FirstName, LastName
FROM SalesLT.Customers
EXCEPT
SELECT FirstName, LastName
FROM SalesLT.Employees;
```

Module 5: Using Functions and Aggregating Data

- Introduction to Built-In Functions
- Scalar Functions
- Aggregate Functions
- Logical Functions
- Window Functions
- Grouping with GROUP BY
- Filtering with HAVING

Introduction to Built-In Functions

Function Category	Description
Scalar	Operate on a single row, return a single value
Aggregate	Take one or more input values, return a single summarizing value
Window	Operate on a window (set) of rows

Scalar Functions

- Operate on elements from a single row as inputs, return a single value as output
- Return a single (scalar) value
- Can be used like an expression in queries
- May be deterministic or non-deterministic

Scalar Function Categories

- Configuration
- Conversion
- Cursor
- Date and Time
- Mathematical
- Metadata
- Security
- String
- System
- System Statistical
- Text and Image

Using Scalar Functions

	SQL Server and Sybase	PostgreSQL
Years	DATEDIFF(yy, start, end)	DATE_PART('YEAR', END) - DATE_PART('YEAR', START)
Months	DATEDIFF(mm, start, end)	Years_diff * 12 + (date_part('month', end) - date_part('month', start))
Days	DATEDIFF(dd, start, end)	DATE_PART('DAY', END - START)
Weeks	DATEDIFF(wk, start, end)	TRUNC(DATE_PART('DAY', END - START)/7)
Hours	DATEDIFF(hh, start, end)	Days_diff * 24 + date_part('hour', end - start)
Minutes	DATEDIFF(mi, start, end)	Hours_diff * 60 + date_part('minute', end - start)
Seconds	DATEDIFF(ss, start, end)	Minutes_diff * 60 + date_part('minute', end - start)
Concat	Concat(string1,string2,..)	CONCAT(STRING1,STRING2,..)

Window Functions

- Functions applied to a window, or set of rows
- Include ranking, offset, aggregate and distribution functions

```
SELECT TOP(3) ProductID, Name, ListPrice,  
             RANK() OVER(ORDER BY ListPrice DESC) AS RankByPrice  
FROM Production.Product  
ORDER BY RankByPrice;
```



ProductID	Name	ListPrice	RankByPrice
8	Gizmo	263.50	1
29	Widget	123.79	2
9	Thingybob	97.00	3

Using Window Functions

```
-- Window functions
SELECT TOP(100) ProductID, Name, ListPrice,
       RANK() OVER(ORDER BY ListPrice DESC) AS RankByPrice
FROM SalesLT.Product AS p
ORDER BY RankByPrice;

SELECT c.Name AS Category, p.Name AS Product, ListPrice,
       RANK() OVER(PARTITION BY c.Name ORDER BY ListPrice DESC) AS RankByPrice
FROM SalesLT.Product AS p
JOIN SalesLT.ProductCategory AS c
ON p.ProductCategoryID = c.ProductcategoryID
ORDER BY Category, RankByPrice;
```

Aggregate Functions

- Functions that operate on sets, or rows of data
- Summarize input rows
- Without GROUP BY clause, all rows are arranged as one group

```
SELECT COUNT(*) AS OrderLines,  
       SUM(OrderQty*UnitPrice) AS TotalSales  
FROM   Sales.OrderDetail;
```



OrderLines	TotalSales
542	71 4002.91 36

Using Aggregate Functions

```
-- Aggregate Functions
```

```
SELECT COUNT(*) AS Products,  
       COUNT(DISTINCT ProductCategoryID) AS Categories,  
       AVG(ListPrice) AS AveragePrice  
FROM SalesLT.Product;
```

```
SELECT COUNT(p.ProductID) BikeModels,  
       AVG(p.ListPrice) AveragePrice  
FROM SalesLT.Product AS p  
JOIN SalesLT.ProductCategory AS c  
ON p.ProductCategoryID = c.ProductCategoryID  
WHERE c.Name LIKE '%Bikes';
```

Grouping with GROUP BY

- GROUP BY creates groups for output rows, according to a unique combination of values specified in the GROUP BY clause
- GROUP BY calculates a summary value for aggregate functions in subsequent phases
- Detail rows are “lost” after GROUP BY clause is processed

```
SELECT CustomerID, COUNT(*) AS Orders  
FROM Sales.SalesOrderHeader  
GROUP BY CustomerID;
```


Grouping with GROUP BY

```
SELECT Salesperson, COUNT(CustomerID) Customers
FROM SalesLT.Customer
GROUP BY Salesperson
ORDER BY Salesperson;
```

```
SELECT c.Salesperson, ISNULL(SUM(oh.SubTotal), 0.00) SalesRevenue
FROM SalesLT.Customer c
LEFT JOIN SalesLT.SalesOrderHeader oh
ON c.CustomerID = oh.CustomerID
GROUP BY c.Salesperson
ORDER BY SalesRevenue DESC;
```

```
SELECT c.Salesperson, CONCAT(c.FirstName + ' ', c.LastName) AS Customer,
       ISNULL(SUM(oh.SubTotal), 0.00) SalesRevenue
FROM SalesLT.Customer c
LEFT JOIN SalesLT.SalesOrderHeader oh
ON c.CustomerID = oh.CustomerID
GROUP BY c.Salesperson, CONCAT(c.FirstName + ' ', c.LastName)
ORDER BY SalesRevenue DESC, Customer;
```

Filtering with HAVING

- HAVING clause provides a search condition that each group must satisfy
- WHERE clause is processed before GROUP BY, HAVING clause is processed after GROUP BY

```
SELECT CustomerID, COUNT(*) AS Orders  
FROM Sales.SalesOrderHeader  
GROUP BY CustomerID  
HAVING COUNT(*) > 10;
```

Filtering with HAVING

```
-- Try to find salespeople with over 150 customers (fails with error)
SELECT Salesperson, COUNT(CustomerID) Customers
FROM SalesLT.Customer
WHERE COUNT(CustomerID) > 100
GROUP BY Salesperson
ORDER BY Salesperson;

--Need to use HAVING clause to filter based on aggregate
SELECT Salesperson, COUNT(CustomerID) Customers
FROM SalesLT.Customer
GROUP BY Salesperson
HAVING COUNT(CustomerID) > 100
ORDER BY Salesperson;
```

Module 6: Using Subquery

- Introduction to Subqueries
- Scalar or Multi-Valued?
- Self-Contained or Correlated?

Introduction to Subqueries

- Subqueries are nested queries: queries within queries
- Results of inner query passed to outer query
 - Inner query acts like an expression from perspective of outer query

SELECT * FROM...



SELECT * FROM...

Scalar or Multi-Valued?

- Scalar subquery returns single value to outer query
 - Can be used anywhere single-valued expression is used: SELECT, WHERE, and so on

```
SELECT orderid, productid, unitprice, qty
FROM Sales.OrderDetails
WHERE orderid =
    (SELECT MAX(orderid) AS lastorder
     FROM Sales.Orders);
```

- Multi-valued subquery returns multiple values as a single column set to the outer query
 - Used with IN predicate

```
SELECT custid, orderid
FROM Sales.orders
WHERE custid IN (
    SELECT custid
    FROM Sales.Customers
    WHERE countryregion = N'Mexico');
```

Self-Contained or Correlated?

- Most subqueries are self-contained and have no connection with the outer query other than passing its results
- Correlated subqueries refer to elements of tables used in outer query
 - Dependent on outer query, cannot be executed separately
 - Behaves as if inner query is executed once per outer row
 - May return scalar value or multiple values

```
SELECT orderid, empid, orderdate
FROM Sales.Orders AS O1
WHERE orderdate = (SELECT MAX(orderdate)
                  FROM Sales.Orders AS O2
                  WHERE O2.empid = O1.empid)
ORDER BY empid, orderdate;
```

Creating a Correlated Subquery

--For each customer list all sales on the last day that they made a sale

```
SELECT CustomerID, SalesOrderID, OrderDate
FROM SalesLT.SalesOrderHeader AS S01
ORDER BY CustomerID, OrderDate
```

```
SELECT CustomerID, SalesOrderID, OrderDate
FROM SalesLT.SalesOrderHeader AS S01
WHERE orderdate =
  (SELECT MAX(orderdate)
   FROM SalesLT.SalesOrderHeader)
```

```
SELECT CustomerID, SalesOrderID, OrderDate
FROM SalesLT.SalesOrderHeader AS S01
WHERE orderdate =
  (SELECT MAX(orderdate)
   FROM SalesLT.SalesOrderHeader AS S02
   WHERE S02.CustomerID = S01.CustomerID)
ORDER BY CustomerID
```



Module 7: Using Table Expression

- Views
- Temporary Tables
- Table Variables
- Table-Valued Functions
- Derived Tables
- Common Table Expressions

Querying Views

- Views are named queries with definitions stored in a database
 - Views can provide abstraction, encapsulation and simplification
 - From an administrative perspective, views can provide a security layer to a database
- Views may be referenced in a SELECT statement just like a table

```
CREATE VIEW Sales.vSalesOrders
AS
SELECT  oh.OrderID, oh.Orderdate, oh.CustomerID,
        od.LineItemNo, od.ProductID, od.Quantity
FROM Sales.OrderHeaders AS oh
JOIN Sales.OrderDetails AS od
ON od.OrderID = oh.OrderID;
```



```
SELECT OrderID, CustomerID, ProductID
FROM Sales.vSalesOrder
ORDER BY OrderID;
```

Temporary Tables

- Temporary tables are used to hold temporary result sets within a user's session
 - Created in schema “pg_temp_” and deleted automatically
 - Created with a CREATE TEMP or CREATE TEMPORARY TABLE

```
CREATE TEMP TABLE tmpProducts
(ProductID INTEGER,
 ProductName varchar(50));
GO
...
SELECT * FROM tmpProducts;
```

Table-Valued Functions

- TVFs are named objects with definitions stored in a database
- TVFs return a virtual table to the calling query
- Unlike views, TVFs support input parameters
 - TVFs may be thought of as parameterized views

```
CREATE OR REPLACE FUNCTION fn_getempinfo(p_employee_id Integer)  
  RETURNS TABLE (employee_id INT, first_name VARCHAR(20),  
                  last_name VARCHAR(25), hire_date date)  
  
AS $$  
BEGIN  
  RETURN QUERY SELECT  
    e.employee_id ,e.first_name ,e.last_name ,e.hire_date  
  FROM  
    public.employees e  
  WHERE e.employee_id = p_employee_id;  
END; $$  
LANGUAGE 'plpgsql';  
  
select * from public.fn_getempinfo(100);
```

Derived Tables – Introduction

- Derived tables are named query expressions created within an outer SELECT statement
- Not stored in database – represents a virtual relational table
- Scope of a derived table is the query in which it is defined

```
SELECT orderyear, COUNT(DISTINCT custid) AS cust_count  
FROM  
    (SELECT YEAR(orderdate) AS orderyear, custid  
     FROM Sales.Orders) AS derived_year  
GROUP BY orderyear;
```

Derived Tables – Guidelines

- Derived tables **must**:
 - Have an alias
 - Have unique names for all columns
 - Not use an ORDER BY clause
 - Not be referred to multiple times in the same query
- Derived tables **may**:
 - Use internal or external aliases for columns
 - Refer to parameters and/or variables
 - Be nested within other derived tables

Derived Tables – Specifying column aliases

- Column aliases may be defined inline:

```
SELECT orderyear, COUNT(DISTINCT custid) AS cust_count  
FROM ( SELECT YEAR(orderdate) AS orderyear, custid  
        FROM Sales.Orders) AS derived_year  
GROUP BY orderyear;
```

- Or externally:

```
SELECT orderyear, COUNT(DISTINCT custid) AS cust_count  
FROM ( SELECT YEAR(orderdate), custid  
        FROM Sales.Orders) AS  
        derived_year(orderyear, custid)  
GROUP BY orderyear;
```

Using Derived Tables

```
SELECT Category, COUNT(ProductID) AS Products
FROM
    (SELECT p.ProductID, p.Name AS Product, c.Name AS Category
     FROM SalesLT.Product AS p
     JOIN SalesLT.ProductCategory AS c
     ON p.ProductCategoryID = c.ProductCategoryID) AS ProdCats
GROUP BY Category
ORDER BY Category;
```


Common Table Expressions (CTEs)

- CTEs are named table expressions defined in a query
- CTEs are similar to derived tables in scope and naming requirements
- Unlike derived tables, CTEs support multiple references and recursion

```
WITH CTE_year (OrderYear, CustID)
AS
(
    SELECT YEAR(orderdate), custid
    FROM Sales.Orders
)
SELECT OrderYear, COUNT(DISTINCT CustID) AS Cust_Count
FROM CTE_year
GROUP BY orderyear;
```

Common Table Expressions – Recursion

- Specify a query for the anchor (root) level
- Use UNION ALL to add a recursive query for other levels
- Query the CTE, with optional MAXRECURSION option

```
WITH OrgReport (ManagerID, EmployeeID, EmployeeName, Level)
AS
(
    SELECT e.ManagerID, e.EmployeeID, EmployeeName, 0
    FROM HR.Employee AS e
    WHERE ManagerID IS NULL
    UNION ALL
    SELECT e.ManagerID, e.EmployeeID, e.EmployeeName, Level + 1
    FROM HR.Employee AS e
    INNER JOIN OrgReport AS o ON e.ManagerID = o.EmployeeID
)
SELECT * FROM OrgReport
OPTION (MAXRECURSION 3);
```

Using Common Table Expressions(CTE)

```
--Using a CTE
WITH ProductsByCategory (ProductID, ProductName, Category)
AS
(
    SELECT p.ProductID, p.Name, c.Name AS Category
    FROM SalesLT.Product AS p
    JOIN SalesLT.ProductCategory AS c
    ON p.ProductCategoryID = c.ProductCategoryID
)

SELECT Category, COUNT(ProductID) AS Products
FROM ProductsByCategory
GROUP BY Category
ORDER BY Category;
```

Using Recursive CTE

```
-- Recursive CTE
SELECT * FROM SalesLT.Employee

-- Using the CTE to perform recursion
WITH OrgReport (ManagerID, EmployeeID, EmployeeName, Level)
AS
(
    -- Anchor query
    SELECT e.ManagerID, e.EmployeeID, EmployeeName, 0
    FROM SalesLT.Employee AS e
    WHERE ManagerID IS NULL

    UNION ALL

    -- Recursive query
    SELECT e.ManagerID, e.EmployeeID, e.EmployeeName, Level + 1
    FROM SalesLT.Employee AS e
    INNER JOIN OrgReport AS o ON e.ManagerID = o.EmployeeID
)

SELECT * FROM OrgReport
OPTION (MAXRECURSION 3);
```

- Thank you

Thank you